

SystemVerilog Hardware Acceleration for Q-Learning

Transcript Date: April 19, 2025

This document captures the technical conversation between a user and ChatGPT regarding the profiling, optimization, and hardware implementation of a Python-based Q-learning algorithm applied to a deterministic FrozenLake environment. The objective was to identify computational bottlenecks in the code and provide equivalent SystemVerilog modules to accelerate the Q-learning process in hardware, suitable for FPGA or ASIC deployment.

1. Repository Analysis

The GitHub repository analyzed is: <https://github.com/ronanmmurphy/Q-Learning-Algorithm>

It implements a Q-learning agent trained to solve a deterministic version of FrozenLake - a 5x5 grid world environment. The agent navigates from the top-left (0,0) to the bottom-right (4,4) while avoiding holes. It employs:

- A State class for managing positions, hole detection, and rewards.
- An Agent class that performs Q-learning with a Q-table structured as nested Python dictionaries.
- Epsilon-greedy policy for exploration-exploitation ($\epsilon = 0.1$).
- Training for 10,000 episodes.

2. Identified Bottlenecks (with Line Numbers)

1. Q-table Dictionary (Lines 110, 176-194): Nested dictionaries ($Q[state][action]$) introduce hashing overhead and poor cache performance. Use fixed-size arrays for fast access.
2. State Object Recreation (Line 166): Re-instantiating State class each episode leads to memory

churn. Reuse one instance with a `reset()` method.

3. Linear Hole Search (Lines 31, 46): Repeated list scans for hole detection can be improved by using a set for constant-time lookup.

4. Random Action Sampling (Line 178): Frequent `np.random.uniform` and `random.choice` calls are relatively slow in tight loops. Consider vectorized pre-generation.

3. Hardware Proposals

3.1 Q-table in SRAM:

- Fixed-size 25x4 Q-table stored in dual-port SRAM.
- Indexed by (state, action) pair using 5-bit and 2-bit indices.
- Read/write supported via constant-time memory access.
- Updates computed in a small ALU: $Q \leftarrow (1 - \alpha) * Q + \alpha * (r + \gamma * \max Q')$.

3.2 State Tracker FSM:

- FSM maintains x and y using 3-bit registers.
- Hole and goal detection via combinational logic or LUT.
- Uses an action decoder (2 bits) to trigger grid movement.
- Reset restores state to (0,0); edge-limited actions check boundaries.

These modules are suitable for synthesis in FPGA or ASIC environments.

4. SystemVerilog Implementations

```
// Q-Table Hardware: SRAM-Based Implementation for 25 states × 4 actions
// Fixed-point 16-bit Q-values, dual-port SRAM for simultaneous read/write

module QTable #(parameter STATE_BITS = 5, // log2(25) = 5
                parameter ACTION_BITS = 2, // log2(4) = 2
                parameter Q_WIDTH = 16)(
    input logic clk,
    input logic rst,
```

```

// Read interface
input logic [STATE_BITS-1:0] read_state,
input logic [ACTION_BITS-1:0] read_action,
output logic [Q_WIDTH-1:0] q_read_data,

// Write interface
input logic we, // write enable
input logic [STATE_BITS-1:0] write_state,
input logic [ACTION_BITS-1:0] write_action,
input logic [Q_WIDTH-1:0] q_write_data
);

// 25 states x 4 actions
localparam ENTRY_COUNT = 25 * 4;
logic [Q_WIDTH-1:0] q_table [0:ENTRY_COUNT-1];

logic [6:0] read_addr, write_addr;
assign read_addr = {read_state, read_action};
assign write_addr = {write_state, write_action};

// Read
assign q_read_data = q_table[read_addr];

// Write
always_ff @(posedge clk) begin
    if (we) begin
        q_table[write_addr] <= q_write_data;
    end
end

endmodule

```

```

// State Tracker FSM for 5x5 Grid (FrozenLake environment)
// Handles (x, y) position, obstacle detection, and reset logic

module StateFSM #(parameter X_BITS = 3, // log2(5) = 3
                  parameter Y_BITS = 3)(
    input logic clk,
    input logic rst,
    input logic [1:0] action, // 2-bit input: 00=up, 01=down, 10=left, 11=right
    input logic step, // step trigger (advance 1 time step)

    output logic [X_BITS-1:0] pos_x,
    output logic [Y_BITS-1:0] pos_y,
    output logic is_hole,
    output logic is_goal
);

// Internal position registers
logic [X_BITS-1:0] x;
logic [Y_BITS-1:0] y;

// Hole ROM: hardcoded 4 positions
function logic is_hole_pos(input logic [2:0] x_in, input logic [2:0] y_in);
    return (x_in == 1 && y_in == 0) ||
           (x_in == 3 && y_in == 1) ||
           (x_in == 4 && y_in == 2) ||
           (x_in == 1 && y_in == 3);
endfunction

function logic is_goal_pos(input logic [2:0] x_in, input logic [2:0] y_in);
    return (x_in == 4 && y_in == 4);
endfunction

// Position update
always_ff @(posedge clk or posedge rst) begin
    if (rst) begin
        x <= 0;
        y <= 0;
    end else if (step) begin
        case (action)
            2'b00: begin
                if (y > 0) y <= y - 1;
                else y <= y;
            end
            2'b01: begin
                if (y < 4) y <= y + 1;
                else y <= y;
            end
            2'b10: begin
                if (x > 0) x <= x - 1;
                else x <= x;
            end
            2'b11: begin
                if (x < 4) x <= x + 1;
                else x <= x;
            end
            default: begin
                x <= x;
                y <= y;
            end
        endcase
    end
end

assign pos_x = x;

```

```
    assign pos_y = y;
    assign is_hole = is_hole_pos(x, y);
    assign is_goal = is_goal_pos(x, y);

endmodule
```